

```
pip install kaggle
```

```
Requirement already satisfied: kaggle in
/opt/anaconda3/lib/python3.11/site-packages (1.7.4.2)
Requirement already satisfied: bleach in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (4.1.0)
Requirement already satisfied: certifi>=14.05.14 in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (2025.1.31)
Requirement already satisfied: charset-normalizer in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (2.0.4)
Requirement already satisfied: idna in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (3.4)
Requirement already satisfied: protobuf in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (4.25.3)
Requirement already satisfied: python-dateutil>=2.5.3 in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (2.8.2)
Requirement already satisfied: python-slugify in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (5.0.2)
Requirement already satisfied: requests in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (2.31.0)
Requirement already satisfied: setuptools>=21.0.0 in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (68.2.2)
Requirement already satisfied: six>=1.10 in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (1.16.0)
Requirement already satisfied: text-unidecode in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (1.3)
Requirement already satisfied: tqdm in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (4.65.0)
Requirement already satisfied: urllib3>=1.15.1 in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (2.0.7)
Requirement already satisfied: webencodings in
/opt/anaconda3/lib/python3.11/site-packages (from kaggle) (0.5.1)
Requirement already satisfied: packaging in
/opt/anaconda3/lib/python3.11/site-packages (from bleach->kaggle)
(23.1)
```

Note: you may need to restart the kernel to use updated packages.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
product_df = pd.read_csv("product_info.csv")
print("Product Info DataFrame loaded. Shape:", product_df.shape)
product_df.head()
```

Product Info DataFrame loaded. Shape: (8494, 27)

```
product_id      product_name  brand_id brand_name
loves_count    \
```

0	P473671	Fragrance Discovery Set	6342	19-69
6320				
1	P473668	La Habana Eau de Parfum	6342	19-69
3827				
2	P473662	Rainbow Bar Eau de Parfum	6342	19-69
3253				
3	P473660	Kasbah Eau de Parfum	6342	19-69
3018				
4	P473658	Purple Haze Eau de Parfum	6342	19-69
2691				

	rating	reviews	size	variation_type
\				
0	3.6364	11.0	NaN	NaN
1	4.1538	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
2	4.2500	16.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
3	4.4762	21.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
4	3.2308	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation

	variation_value	...	online_only	out_of_stock	sephora_exclusive	\
0	NaN	...	1	0	0	
1	3.4 oz/ 100 mL	...	1	0	0	
2	3.4 oz/ 100 mL	...	1	0	0	
3	3.4 oz/ 100 mL	...	1	0	0	
4	3.4 oz/ 100 mL	...	1	0	0	

	highlights	primary_category
\		
0	['Unisex/ Genderless Scent', 'Warm &Spicy Scen...	Fragrance
1	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance
2	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance
3	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance
4	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance

	secondary_category	tertiary_category	child_count	child_max_price
\				
0	Value & Gift Sets	Perfume Gift Sets	0	NaN
1	Women	Perfume	2	85.0
2	Women	Perfume	2	75.0

3	Women	Perfume	2	75.0
4	Women	Perfume	2	75.0

	child_min_price
0	NaN
1	30.0
2	30.0
3	30.0
4	30.0

[5 rows x 27 columns]

```
import os
print(os.getcwd())
```

/Users/alexandriapetersen/Downloads

```
import glob
```

```
csv_files = glob.glob("*.csv")
print(csv_files)
```

```
['product_info.csv', 'reviews_0-250.csv', 'reviews_1250-end.csv',
'reviews_750-1250.csv', 'image_labels.csv', 'target.csv',
'LaptopSales2008.csv', 'reviews_250-500.csv', 'reviews_500-750.csv']
```

```
import os
print(os.getcwd())
```

/Users/alexandriapetersen/Downloads

```
import os
```

```
os.chdir("/Users/alexandriapetersen/Downloads")
print("Now in:", os.getcwd())
```

Now in: /Users/alexandriapetersen/Downloads

```
product_df = pd.read_csv("product_info.csv")
print("Product Info DataFrame loaded. Shape:", product_df.shape)
product_df.head()
```

Product Info DataFrame loaded. Shape: (8494, 27)

	product_id	product_name	brand_id	brand_name
loves_count \				
0	P473671	Fragrance Discovery Set	6342	19-69
6320				
1	P473668	La Habana Eau de Parfum	6342	19-69

3827

2 P473662 Rainbow Bar Eau de Parfum 6342 19-69

3253

3 P473660 Kasbah Eau de Parfum 6342 19-69

3018

4 P473658 Purple Haze Eau de Parfum 6342 19-69

2691

	rating	reviews	size	variation_type
\				
0	3.6364	11.0	NaN	NaN

1	4.1538	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
---	--------	------	----------------	------------------------------------

2	4.2500	16.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
---	--------	------	----------------	------------------------------------

3	4.4762	21.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
---	--------	------	----------------	------------------------------------

4	3.2308	13.0	3.4 oz/ 100 mL	Size + Concentration + Formulation
---	--------	------	----------------	------------------------------------

	variation_value	... online_only	out_of_stock	sephora_exclusive	\
0	NaN	...	1	0	0
1	3.4 oz/ 100 mL	...	1	0	0
2	3.4 oz/ 100 mL	...	1	0	0
3	3.4 oz/ 100 mL	...	1	0	0
4	3.4 oz/ 100 mL	...	1	0	0

	highlights	primary_category
\		
0	['Unisex/ Genderless Scent', 'Warm &Spicy Scen...	Fragrance
1	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance
2	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance
3	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance
4	['Unisex/ Genderless Scent', 'Layerable Scent'...	Fragrance

	secondary_category	tertiary_category	child_count	child_max_price
\				
0	Value & Gift Sets	Perfume Gift Sets	0	NaN
1	Women	Perfume	2	85.0
2	Women	Perfume	2	75.0
3	Women	Perfume	2	75.0

4	Women	Perfume	2	75.0
---	-------	---------	---	------

	child_min_price
0	NaN
1	30.0
2	30.0
3	30.0
4	30.0

[5 rows x 27 columns]

```
import pandas as pd
import glob
```

```
# Load all review files
```

```
review_files = glob.glob("reviews_*.csv")
```

```
# Combine into one DataFrame
```

```
reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
ignore_index=True)
```

```
print("Reviews DataFrame loaded. Shape:", reviews_df.shape)
```

```
reviews_df.head()
```

```
/var/folders/w3/4bkcyhf15_3981_dm0cmtps80000gp/T/
```

```
ipykernel_61923/3215166159.py:8: DtypeWarning: Columns (1) have mixed
types. Specify dtype option on import or set low_memory=False.
```

```
reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
ignore_index=True)
```

```
/var/folders/w3/4bkcyhf15_3981_dm0cmtps80000gp/T/ipykernel_61923/32151
```

```
66159.py:8: DtypeWarning: Columns (1) have mixed types. Specify dtype
option on import or set low_memory=False.
```

```
reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
ignore_index=True)
```

```
/var/folders/w3/4bkcyhf15_3981_dm0cmtps80000gp/T/ipykernel_61923/32151
```

```
66159.py:8: DtypeWarning: Columns (1) have mixed types. Specify dtype
option on import or set low_memory=False.
```

```
reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
ignore_index=True)
```

```
Reviews DataFrame loaded. Shape: (1094411, 19)
```

	Unnamed: 0	author_id	rating	is_recommended	helpfulness	\
0	0	1741593524	5	1.0	1.0	
1	1	31423088263	1	0.0	NaN	
2	2	5061282401	5	1.0	NaN	
3	3	6083038851	5	1.0	NaN	
4	4	47056667835	5	1.0	NaN	

total_feedback_count	total_neg_feedback_count
----------------------	--------------------------

	total_pos_feedback_count	\
0	2	0
2		
1	0	0
0		
2	0	0
0		
3	0	0
0		
4	0	0
0		

	submission_time	review_text	\
0	2023-02-01	I use this with the Nudestix "Citrus Clean Bal...	
1	2023-03-21	I bought this lip mask after reading the revie...	
2	2023-03-21	My review title says it all! I get so excited ...	
3	2023-03-20	I've always loved this formula for a long time...	
4	2023-03-20	If you have dry cracked lips, this is a must h...	

	review_title	skin_tone	eye_color
skin_type \			
0	Taught me how to double cleanse!	NaN	brown dry
1	Disappointed	NaN	NaN NaN
2	New Favorite Routine	light	brown dry
3	Can't go wrong with any of them	NaN	brown combination
4	A must have !!!	light	hazel combination

	hair_color	product_id
product_name \		
0	black	P504322 Gentle Hydra-Gel Face Cleanser
1	NaN	P420652 Lip Sleeping Mask Intense Hydration with Vitam...
2	blonde	P420652 Lip Sleeping Mask Intense Hydration with Vitam...
3	black	P420652 Lip Sleeping Mask Intense Hydration with Vitam...
4	NaN	P420652 Lip Sleeping Mask Intense Hydration with Vitam...

	brand_name	price_usd
0	NUDESTIX	19.0
1	LANEIGE	24.0
2	LANEIGE	24.0
3	LANEIGE	24.0
4	LANEIGE	24.0

```
print("Reviews columns:\n", reviews_df.columns)
print("\nProduct columns:\n", product_df.columns)
```

```
Reviews columns:
Index(['Unnamed: 0', 'author_id', 'rating', 'is_recommended',
      'helpfulness',
      'total_feedback_count', 'total_neg_feedback_count',
      'total_pos_feedback_count', 'submission_time', 'review_text',
      'review_title', 'skin_tone', 'eye_color', 'skin_type',
      'hair_color',
      'product_id', 'product_name', 'brand_name', 'price_usd'],
      dtype='object')
```

```
Product columns:
Index(['product_id', 'product_name', 'brand_id', 'brand_name',
      'loves_count',
      'rating', 'reviews', 'size', 'variation_type',
      'variation_value',
      'variation_desc', 'ingredients', 'price_usd',
      'value_price_usd',
      'sale_price_usd', 'limited_edition', 'new', 'online_only',
      'out_of_stock', 'sephora_exclusive', 'highlights',
      'primary_category',
      'secondary_category', 'tertiary_category', 'child_count',
      'child_max_price', 'child_min_price'],
      dtype='object')
```

```
# Merge on 'product_id'
merged_df = reviews_df.merge(product_df, on='product_id', how='left')
```

```
# Check results
```

```
print("Merged DataFrame shape:", merged_df.shape)
merged_df.head()
```

```
Merged DataFrame shape: (1094411, 45)
```

	Unnamed: 0	author_id	rating_x	is_recommended	helpfulness	\
0	0	1741593524	5	1.0	1.0	
1	1	31423088263	1	0.0	NaN	
2	2	5061282401	5	1.0	NaN	
3	3	6083038851	5	1.0	NaN	
4	4	47056667835	5	1.0	NaN	

total_feedback_count	total_neg_feedback_count
----------------------	--------------------------

	total_pos_feedback_count	\
0	2	0
2		
1	0	0
0		
2	0	0
0		
3	0	0
0		
4	0	0
0		

	submission_time	review_text	...	\
0	2023-02-01	I use this with the Nudestix "Citrus Clean	Bal...	...
1	2023-03-21	I bought this lip mask after reading the	revie...	...
2	2023-03-21	My review title says it all! I get so	excited	...
3	2023-03-20	I've always loved this formula for a long	time...	...
4	2023-03-20	If you have dry cracked lips, this is a must	h...	...

	online_only	out_of_stock	sephora_exclusive	\
0	1	0	0	
1	0	0	1	
2	0	0	1	
3	0	0	1	
4	0	0	1	

	highlights	primary_category
0	['Clean at Sephora']	Skincare
1	['allure 2019 Best of Beauty Award Winner', 'C...]	Skincare
2	['allure 2019 Best of Beauty Award Winner', 'C...]	Skincare
3	['allure 2019 Best of Beauty Award Winner', 'C...]	Skincare
4	['allure 2019 Best of Beauty Award Winner', 'C...]	Skincare

	secondary_category	tertiary_category	child_count
0	Cleansers	NaN	0
1	Lip Balms & Treatments	NaN	3


```

24.0
2 Lip Balms & Treatments      NaN      3
24.0
3 Lip Balms & Treatments      NaN      3
24.0
4 Lip Balms & Treatments      NaN      3
24.0

```

```

child_min_price
0      NaN
1     24.0
2     24.0
3     24.0
4     24.0

```

```
[5 rows x 45 columns]
```

```

missing = merged_df[merged_df['ingredients'].isnull()]
print("Reviews with missing product info:", missing.shape[0])
Reviews with missing product info: 22025

```

RATINGS DISTRIBUTION

```

print("Columns in merged_df:")
print(merged_df.columns)

Columns in merged_df:
Index(['Unnamed: 0', 'author_id', 'rating_x', 'is_recommended',
      'helpfulness',
      'total_feedback_count', 'total_neg_feedback_count',
      'total_pos_feedback_count', 'submission_time', 'review_text',
      'review_title', 'skin_tone', 'eye_color', 'skin_type',
      'hair_color',
      'product_id', 'product_name_x', 'brand_name_x', 'price_usd_x',
      'product_name_y', 'brand_id', 'brand_name_y', 'loves_count',
      'rating_y',
      'reviews', 'size', 'variation_type', 'variation_value',
      'variation_desc', 'ingredients', 'price_usd_y',
      'value_price_usd',
      'sale_price_usd', 'limited_edition', 'new', 'online_only',
      'out_of_stock', 'sephora_exclusive', 'highlights',
      'primary_category',
      'secondary_category', 'tertiary_category', 'child_count',
      'child_max_price', 'child_min_price'],
      dtype='object')

print(merged_df.columns)


```

```
Index(['Unnamed: 0', 'author_id', 'rating_x', 'is_recommended',
      'helpfulness',
      'total_feedback_count', 'total_neg_feedback_count',
      'total_pos_feedback_count', 'submission_time', 'review_text',
      'review_title', 'skin_tone', 'eye_color', 'skin_type',
      'hair_color',
      'product_id', 'product_name_x', 'brand_name_x', 'price_usd_x',
      'product_name_y', 'brand_id', 'brand_name_y', 'loves_count',
      'rating_y',
      'reviews', 'size', 'variation_type', 'variation_value',
      'variation_desc', 'ingredients', 'price_usd_y',
      'value_price_usd',
      'sale_price_usd', 'limited_edition', 'new', 'online_only',
      'out_of_stock', 'sephora_exclusive', 'highlights',
      'primary_category',
      'secondary_category', 'tertiary_category', 'child_count',
      'child_max_price', 'child_min_price'],
      dtype='object')
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
plt.figure(figsize=(8, 5))
sns.countplot(x=merged_df['rating_x'].dropna(), palette='viridis')
```

```
plt.title('Distribution of Product Ratings (from Reviews)')
plt.xlabel('Rating')
plt.ylabel('Number of Reviews')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
merged_df.rename(columns={'rating_x': 'rating'}, inplace=True)
```

Break ratings down by skin type (e.g., oily vs dry)

```
avg_ratings_by_skin = merged_df.groupby('skin_type')
['rating'].mean().sort_values(ascending=False)
print(avg_ratings_by_skin)

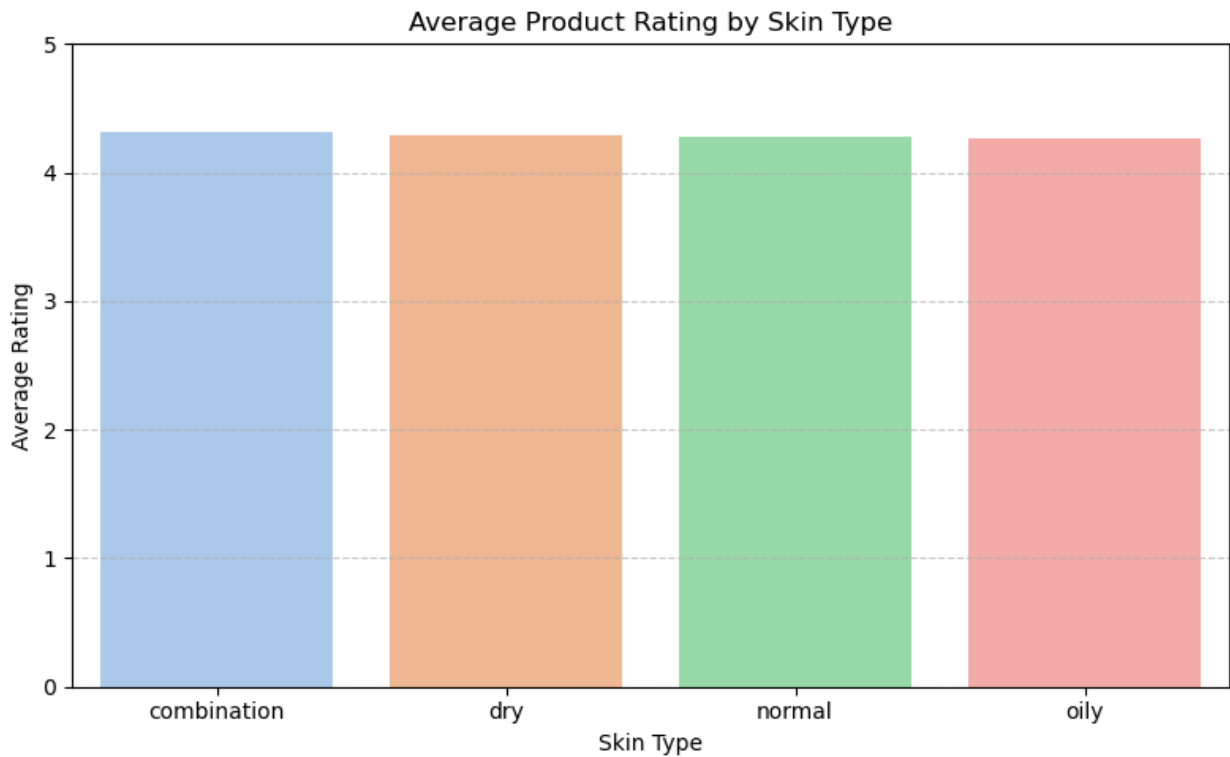
skin_type
combination    4.309339
dry            4.291249
normal         4.282276
oily           4.270910
Name: rating, dtype: float64

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8, 5))
sns.barplot(x=avg_ratings_by_skin.index, y=avg_ratings_by_skin.values,
palette='pastel')

plt.title('Average Product Rating by Skin Type')
plt.xlabel('Skin Type')
plt.ylabel('Average Rating')
plt.ylim(0, 5)
```

```
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
skin_type_counts = merged_df['skin_type'].value_counts()
print(skin_type_counts)
```

```
skin_type
combination    544513
dry             185937
normal         131910
oily           120494
Name: count, dtype: int64
```

Higher average = more satisfaction for that skin type Large differences might indicate some products work better for oily vs dry vs combination skin

Interpretation of Your Plot: Your bar chart shows average product rating by skin type. Based on what I see:


All skin types have similar average ratings, hovering slightly above 4.3. Combination skin seems to rate products the highest on average. Normal, oily, and dry are close behind — there isn't a drastic difference, but there may still be nuances worth exploring.

□ Number of Reviews by Skin Type

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(7, 5))
sns.countplot(data=merged_df, x='skin_type',
order=merged_df['skin_type'].value_counts().index, palette='Set2')

plt.title('Number of Reviews by Skin Type')
plt.xlabel('Skin Type')
plt.ylabel('Number of Reviews')
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



Why this matters: If 70% of data is from people with oily skin, then high average ratings might just reflect their preferences. Equal or balanced review counts across skin types = more reliable comparisons.

Interpretation: Number of Reviews by Skin Type

1. Combination skin dominates the dataset There are over 500,000 reviews from users with combination skin. That's more than half the dataset, meaning this skin type is

heavily overrepresented. This could skew overall average ratings upward if products tend to work better for this group.

2. Dry, normal, and oily skin types are underrepresented Dry skin: 180,000 reviews Normal skin: 130,000 reviews Oily skin: 120,000 reviews These are much smaller sample sizes, so we need to be cautious when interpreting average ratings for these groups.

Why This Matters: The average rating for combination skin might be more statistically reliable than for other skin types. If you're building a prediction model, you'll want to: Either balance the classes (resample or reweight) Or take this imbalance into account when interpreting results

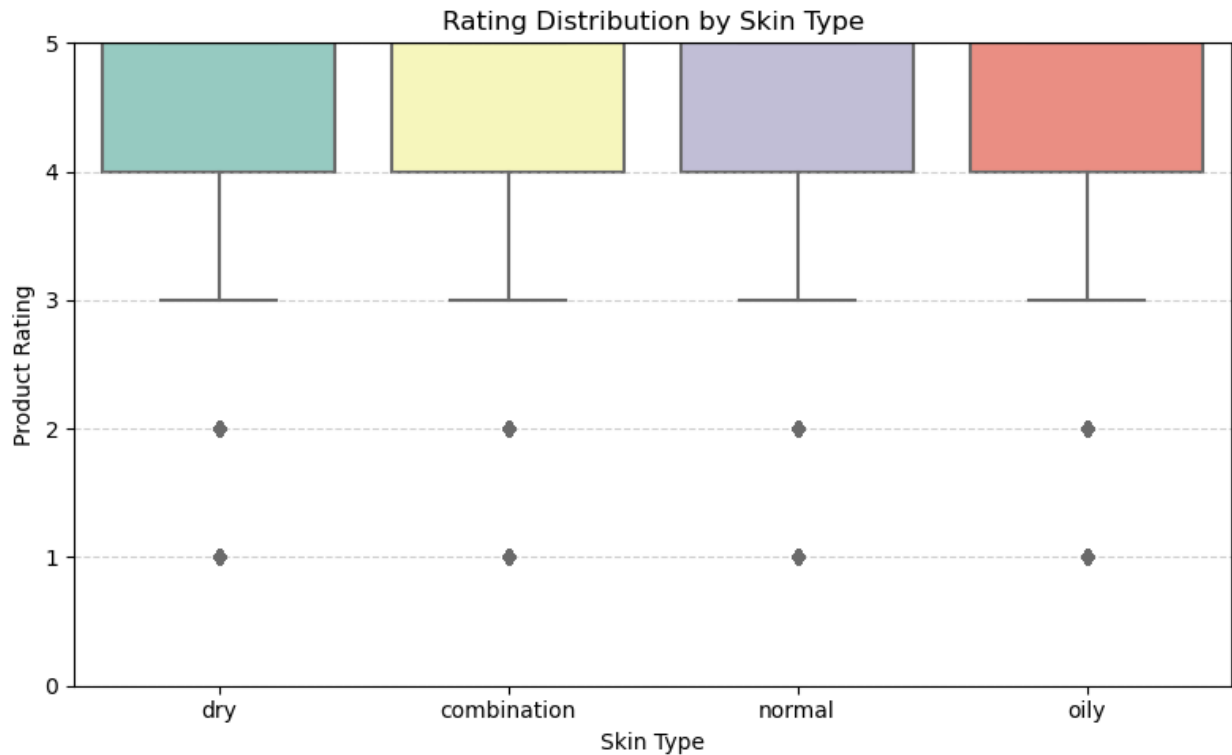
What You Can Do Next: Use boxplots to see how rating variability differs across skin types Look at top-rated products per skin type Model product effectiveness (rating ≥ 4) using skin type as a feature

Plot Rating Distribution by Skin Type (Boxplot)

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))
sns.boxplot(data=merged_df, x='skin_type', y='rating', palette='Set3')

plt.title('Rating Distribution by Skin Type')
plt.xlabel('Skin Type')
plt.ylabel('Product Rating')
plt.ylim(0, 5)
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



What This Tells Us: High overall satisfaction: Skincare products on Sephora are generally well-rated. Slight limitations: With everyone rating so high, it may be hard to find meaningful variation without deeper features like ingredients or specific product types. Good baseline: Skin type may not dramatically affect average satisfaction—but it may still interact with certain ingredients or product categories.

Logistic Regression Predict whether a product will get a positive rating (e.g., rating ≥ 4) Inputs (Features): Skin type (one-hot encoded) Ingredients (presence/absence or TF-IDF vectorized) Brand (optional) Price (optional)

```
merged_df['positive_rating'] = (merged_df['rating'] >= 4).astype(int)
y = merged_df['positive_rating']

from sklearn.preprocessing import OneHotEncoder

# One-hot encode skin_type
skin_dummies = pd.get_dummies(merged_df['skin_type'], prefix='skin')

# Add numeric features
X = pd.concat([skin_dummies, merged_df[['price_usd_x']]], axis=1)

# Drop rows with missing price
X = X.dropna()
y = y.loc[X.index] # Align y to X

from sklearn.model_selection import train_test_split
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

# Evaluate
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

Accuracy: 0.8199586080234645
Confusion Matrix:
[[ 0 39408]
 [ 0 179475]]

```

Model Results Explained Accuracy: 81.99% That looks great at first glance — over 81% of predictions were correct. What this means: The model predicted "positive rating" for every review It never predicted a negative review This is likely because the data is imbalanced (way more 4–5 star ratings than 1–3)

Here's How To Improve It

1. Check Class Balance:

```

y.value_counts(normalize=True)

positive_rating
1    0.820843
0    0.179157
Name: proportion, dtype: float64

# Combine X and y
full_data = pd.concat([X, y], axis=1)

# Separate positive and negative classes
pos = full_data[full_data['positive_rating'] == 1]
neg = full_data[full_data['positive_rating'] == 0]

# Undersample positive class to match negatives
pos_downsampled = pos.sample(len(neg), random_state=42)

# Combine for balanced dataset
balanced = pd.concat([neg, pos_downsampled])

# Split

```



```

X_balanced = balanced.drop('positive_rating', axis=1)
y_balanced = balanced['positive_rating']

X_train, X_test, y_train, y_test = train_test_split(X_balanced,
y_balanced, test_size=0.2, random_state=42)

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

Accuracy: 0.5113669688508078
Confusion Matrix:
[[19992 19347]
 [18976 20114]]

```

Accuracy: 51.1% Lower than before (because now dealing with a balanced dataset) But much more realistic and meaningful Accuracy around 50% is expected for a 50/50 class split if the model is still learning

What This Means: 20k true positives (1 predicted as 1) 20k true negatives (0 predicted as 0) Model is doing okay, but still making a lot of mistakes (especially false positives and false negatives)

Next Steps to Improve the Model

1. Try a Better Classifier: Random Forest This is a powerful model that handles:

Nonlinear patterns Feature interactions Imbalanced classes (much better than logistic regression)

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test,
y_pred))

Accuracy: 0.5463157760522256
Confusion Matrix:
[[21266 18073]
 [17509 21581]]


Classification Report:
              precision    recall  f1-score   support


```

0	0.55	0.54	0.54	39339
1	0.54	0.55	0.55	39090
accuracy			0.55	78429
macro avg	0.55	0.55	0.55	78429
weighted avg	0.55	0.55	0.55	78429

```
import matplotlib.pyplot as plt

feature_importance = pd.Series(rf.feature_importances_,
index=X.columns)
feature_importance.nlargest(10).plot(kind='barh')
plt.title("Top 10 Most Important Features")
plt.show()
```



Product price is the strongest predictor of positive skincare reviews—more than skin type. Label your x-axis: Feature Importance Label y-axis: Model Input Features Color code bar for price_usd_x to highlight it Interpretation: Price (price_usd_x) had the highest impact on predicting if a product gets a high rating Skin types (skin_oily, skin_dry, etc.) had very little impact on model predictions Suggests consumer perception of quality may be influenced more by product cost than personalization Implications: Brands may benefit from strategic pricing or value messaging Skin-type targeting alone may not drive higher satisfaction Recommender systems may be improved by incorporating ingredients and user preferences, not just skin type

Let's build a data-backed recommendation rule like:

"For dry skin, recommend products with [X ingredients] and a [mid-range price]."

```
dry_df = merged_df[merged_df['skin_type'] == 'dry']

top_dry_products = (
    dry_df.groupby('product_name_x')
        .agg(avg_rating=('rating', 'mean'), review_count=('rating',
        'count'), avg_price=('price_usd_x', 'mean'))
        .query('review_count >= 50') # filter for reliability
        .sort_values(by='avg_rating', ascending=False)
        .head(10)
)

top_dry_products.reset_index(inplace=True)
top_dry_products
```

	product_name_x	avg_rating	\
0	Barrier+ Triple Lipid + Collagen + Niacinamide...	4.888889	
1	Super Rich Repair Moisturizer	4.888889	
2	Evercalm Barrier Support Face Oil	4.884615	
3	Luxury Sun Ritual Pore Smoothing Sunscreen SPF 30	4.881356	
4	Ultralight Moisture-Boosting Botanical Oil	4.873239	
5	Truth Barrier Booster Orange Ferment Vitamin C...	4.866071	
6	Silk Rice Makeup-Removing Cleansing Oil	4.858974	
7	Rénergie Lift Multi-Action Ultra Dark Spot Cor...	4.850000	
8	Daily Milkfoliant Exfoliator	4.847328	
9	Juneberry & Collagen Hydrating Cold Cream Clea...	4.843137	

	review_count	avg_price
0	117	69.0
1	54	94.0
2	52	60.0
3	59	38.0
4	71	44.0
5	112	48.0
6	78	46.0
7	60	135.0
8	131	65.0
9	51	39.0

```
ingredient_lookup = merged_df[['product_name_x',
'ingredients']].drop_duplicates()

# Merge to get ingredients for top products
top_dry_products = top_dry_products.merge(ingredient_lookup,
on='product_name_x', how='left')
```

```
for i, row in top_dry_products.iterrows():
    print(f"\nProduct: {row['product_name_x']}\nPrice: $
{row['avg_price']:.2f}\nIngredients:\n{row['ingredients']}")
```

Product: Barrier+ Triple Lipid + Collagen + Niacinamide Activating Serum

Price: \$69.00

Ingredients:

['Water/Aqua/Eau, Caprylyl Caprylate/Caprates, Propanediol, Jojoba Oil/Macadamia Seed Oil Esters, Niacinamide, Glycerin, Ammonium Acryloyldimethyltaurate/VP Copolymer, Squalene, Macadamia Integrifolia Seed Oil, Amylopectin, Lecithin, Phytosteryl Macadamiate, Collagen Amino Acids, Lithothamnion Calcareum Extract, Sodium Hyaluronate, Oligopeptide-3, Oligopeptide-2, Oligopeptide-1, Hexapeptide-11, Folic Acid, Bacillus/Soybean Ferment Extract, Lactic Acid, Phospholipids, Ceramide NP, Phytosterols, Phytosphingosine, Ceramide AP, Cholesterol, Myrica Cerifera (Bayberry) Fruit Extract, Akebia Quinata Stem Extract, Prunus Lannesiana (Cherry Blossom) Flower Extract, Saccharomyces Lysate, Lactobacillus Ferment Lysate, Tripeptide-1, Ceramide EOP, Polyglutamic Acid, Tocopheryl Acetate, Tocopherol, Butylene Glycol, Caprylyl Glycol, Sodium Lauroyl Lactylate, Ethylhexylglycerin, Hexylene Glycol, Pentyleneglycol, Acetyl Glutamine, Trisodium Ethylenediamine Disuccinate, Leuconostoc/Radish Root Ferment Filtrate, Sodium Benzoate, Hydroxyacetophenone, Xanthan Gum, Potassium Sorbate, Carbomer, 1,2-Hexanediol, Dextran, Phenoxyethanol.']

Product: Super Rich Repair Moisturizer

Price: \$94.00

Ingredients:

['Water/Aqua/Eau, Isohexadecane, Dipropylene Glycol, Caprylic/Capric/Myristic/Stearic Triglyceride, Lauryl PEG-9 Polydimethylsiloxylethyl Dimethicone, Simmondsia Chinensis (Jojoba) Seed Oil, Butyrospermum Parkii (Shea) Butter, Dimethicone, Glycerin, Octyldodecyl Neopentanoate, Sodium Chloride, Isostearic Acid, Arginine/Lysine Polypeptide, Palmitoyl Tripeptide-5, Avena Sativa (Oat) Kernel Extract, Pyrus Malus (Apple) Seed Extract, Borago Officinalis Seed Oil, Oenothera Biennis (Evening Primrose) Oil, Hydrogenated Coconut Oil, Gardenia Taitensis Flower, Laminaria Digitata Extract, Yeast Extract, Glucosamine HCL, Cedrus Atlantica Bark Oil, Cupressus Sempervirens Leaf/Stem Extract, Eucalyptus Globulus Leaf Oil, Helianthus Annuus (Sunflower) Seed Oil, Pelargonium Graveolens Flower Oil, Aniba Rosodora (Rosewood) Wood Oil, Abies Sibirica Oil, Santalum Album (Sandalwood) Oil, Sodium Hyaluronate, Tocopheryl Acetate, Colloidal Oatmeal, Madecassoside, Urea, Hexyldecylol, Polyglyceryl-4 Isostearate, Cetyl Dimethicone, Hydrogenated Castor Oil, Butylene Glycol, Caprylyl Glycol, Ethylhexylglycerin, Hexylene Glycol, Linalool, Geraniol, Citronellol, Phenoxyethanol.']

Product: Evercalm Barrier Support Face Oil

Price: \$60.00

Ingredients:

['Caprylic/Capric Triglyceride, Oryza Sativa Bran Oil, Oryza Sativa Germ Oil, Camelina Sativa Seed Oil, Plukenetia Volubilis Seed Oil, Camellia Japonica Seed Oil, Limnanthes Alba Seed Oil, Rosa Canina Fruit Oil, Bisabolol, Tocopherol.']

Product: Luxury Sun Ritual Pore Smoothing Sunscreen SPF 30

Price: \$38.00

Ingredients:

['Zinc Oxide 10%, Water, Butylene Glycol, Caprylic/Capric Triglyceride, Dimethicone, Cetearyl Alcohol, Cetearyl Oliviate, Polyglyceryl-3 Diisostearate, Sorbitan Oliviate, Jasminum Officinale (Jasmine) Extract, Hibiscus Sabdariffa Flower Extract, Adenium Obesum (Desert Rose) Extract, Cetearyl Glucoside, Triethoxycaprylylsilane, Camellia Sinensis (Green Tea) Leaf Oil, Tocopheryl Acetate, Xanthan Gum, Phenoxyethanol, Ethylhexylglycerin, Iron Oxides.']

Product: Ultralight Moisture-Boosting Botanical Oil

Price: \$44.00

Ingredients:

['Caprylic/Capric Triglyceride, Squalane, Hydrogenated Ethylhexyl Oliviate, Diheptyl Succinate, Simmondsia Chinensis (Jojoba) Seed Oil, Crambe Abyssinica Seed Oil, Euterpe Oleracea Fruit Oil, Helianthus Annuus (Sunflower) Seed Oil, Vitis Vinifera (Grape) Seed Oil, Sclerocarya Birrea Seed Oil, Persea Gratissima (Avocado) Oil, Cassia Angustifolia Seed Polysaccharide, Argania Spinosa Kernel Oil, Brassica Campestris (Rapeseed) Seed Oil, Moringa Oleifera Seed Oil, Hippophae Rhamnoides Oil, Linum Usitatissimum (Linseed) Seed Oil, Opuntia Ficus-Indica Seed Oil, Sesamum Indicum (Sesame) Seed Oil, Emblica Officinalis Fruit Extract, Capryloyl Glycerin/Sebacic Acid Copolymer, Hydrogenated Olive Oil Unsaponifiables, Water (Aqua, Eau).']

Product: Truth Barrier Booster Orange Ferment Vitamin C Essence

Price: \$48.00

Ingredients:

['Aqua/water/eau, Lactobacillus Ferment, 3-o-ethyl Ascorbic Acid, Niacinamide, Hyaluronic Acid, Sodium Hyaluronate, Panthenol, Sodium Polyglutamate, Tocopherol, Glycerin, Citrus Aurantium Dulcis (Orange) Peel Extract, Citrus Aurantium Dulcis (Orange) Callus Culture Extract, Citrus Sinensis (Orange) Fruit Extract, Citrus Aurantium Dulcis (Orange) Fruit Extract, Citrus Aurantium Dulcis (Orange) Oil, Hippophae Rhamnoides Extract, Lycium Barbarum Fruit Extract, Rosa Canina Fruit Extract, Citrus Limon (Lemon) Fruit Extract, Gluconolactone, Lithothamnion Calcareum Extract, Helianthus Annuus (Sunflower) Seed Oil, Chondrus Crispus (Carrageenan) Extract, Leuconostoc/radish Root Ferment Filtrate, Citric Acid, Sodium Riboflavin Phosphate, Sodium Tocopheryl Phosphate, Sodium Citrate, Potassium Hydroxide, Xanthan Gum, Sodium Benzoate, Calcium Gluconate,

Triethylhexanoin, Polyglyceryl-3 Laurate, Acrylates/c10-30 Alkyl Acrylate Crosspolymer, Pentaerythrityl Tetra-di-t-butyl Hydroxyhydrocinnamate, Phenoxyethanol, Chlorphenesin, Triethyl Citrate, Limonene, Linalool, Citral.']

Product: Silk Rice Makeup-Removing Cleansing Oil

Price: \$46.00

Ingredients:

['Caprylic/Capric Triglyceride, Octyldodecanol, Dicaprylyl Ether, Polyglyceryl-10 Diisostearate, Oryza Sativa (Rice) Bran Oil, Polyglyceryl-4 Caprate, Water, Tocopherol, Helianthus Annuus (Sunflower) Seed Oil.']

Product: Rénergie Lift Multi-Action Ultra Dark Spot Correcting Cream SPF 30

Price: \$135.00

Ingredients:

['Avobenzene 3%, Octisalate 5%, Octocrylene 7%, Water, Glycerin, Dimethicone, Isononyl Isononanoate, Propanediol, Vinyl Dimethicone/Methicone Silsesquioxane Crosspolymer, Alcohol Denat., Bis-Peg-18 Methyl Ether Dimethyl Silane, Polyglyceryl-6 Distearate, Jojoba Esters, Tocopherol, Limnanthes Alba (Meadowfoam) Seed Oil, Acacia Decurrens Flower Wax, Guanosine, Cyathea Medullaris Leaf Extract, Sodium Hyaluronate, hydrolyzed linseed extract, Sodium Hydroxide, Sodium Dodecylbenzenesulfonate, Sodium Benzoate, Red 33, Sodium Levulinate, Phenoxyethanol, Adenosine, Peg-8 Laurate, Helianthus Annuus (Sunflower) Seed Wax, Polyglyceryl-3 Beeswax, Polyglycerin-3, Ammonium Acryloyldimethyltaurate/Steareth-25 Methacrylate Crosspolymer, Dimethicone/Vinyl Dimethicone Crosspolymer, Dimethiconol, Limonene, Xanthan Gum, Benzyl Alcohol, Cinnamic Acid, Leontopodium Alpinum Flower/Leaf Extract, Capryloyl Salicylic Acid, Caprylyl Glycol, Geraniol, Disodium Stearoyl Glutamate, Disodium Edta, Cetyl Alcohol, Citric Acid, Potassium Sorbate, Scutellaria Baicalensis Root Extract, Levulinic Acid, Styrene/Acrylates Copolymer, Glyceryl Caprylate, Fragrance.']

Product: Daily Milkfoliant Exfoliator

Price: \$65.00

Ingredients:

['Sodium CocoylIsethionate, Microcrystalline Cellulose, Sodium Bicarbonate, Sorbitol, ZeaMays (Corn) Starch, Sodium CocoylGlutamate, Saccharomyces Ferment, Oryza Sativa (Rice) Starch, Magnesium Oxide, Citric Acid, Tapioca Starch, AvenaSativa (Oat) Kernel Protein, Camellia Oleifera Seed Extract, Silica, Cocos Nucifera (Coconut) Fruit Extract, Inulin, Cassia HydroxypropyltrimoniumChloride, Silybum Marianum Seed Oil, Kaolin, Lauryl Methacrylate/Glycol DimethacrylateCrosspolymer, Papain, Hyaluronic Acid, AvenaSativa (Oat) Kernel Extract, Helianthus Annuus (Sunflower) Seed Oil, AvenaSativa (Oat) Bran Extract, Glycine Soja (Soybean) Oil, Citrus Aurantium Dulcis (Orange) Peel Extract, Vitis Vinifera (Grape) Fruit Extract,

Citrus Aurantium Bergamia (Bergamot) Fruit Oil, Aniba Rosodora (Rosewood) Wood Oil, Salvia Sclarea (Clary) Oil, Eucalyptus Globulus Leaf Oil, Chamomilla Recutita (Matricaria) Flower Oil, Cucurbita Pepo (Pumpkin) Fruit Extract, Cucumis Sativus (Cucumber) Fruit Extract, Perilla Ocymoides Leaf Extract, Daucus Carota Sativa (Carrot) Root Extract, Sea Salt (Maris Sal), Juniperus Virginiana Oil, Alpha-Glucan Oligosaccharide, Tocopherol, Castoryl Maleate, Caprylyl Glycol, Allantoin, Diglycerin, Maltodextrin, Lauroyl Lysine, Arginine, Saccharide Isomerate, Sodium Lauroyl Glutamate, Beta-Carotene, Xanthan Gum, Ethylhexylglycerin, 1,2-Hexanediol, Hydroxypropyl Methylcellulose, Dextrin, Sodium Palmitate, Sodium Chloride, Acacia Senegal Gum, Sodium Citrate, Water/Aqua/Eau, 0-Cymen-5-ol, Limonene, Linalool.']

Product: Juneberry & Collagen Hydrating Cold Cream Cleanser

Price: \$39.00

Ingredients:

['Water/Aqua/Eau, Helianthus Annuus (Sunflower) Seed Oil, Cetyl Alcohol, Glycerin, Cetearyl Alcohol, Cetearyl Olivatate, Sorbitan Olivatate, Theobroma Cacao (Cocoa) Seed Butter*, Propanediol, Kaolin, Decyl Glucoside, Bakuchiol, Collagen, Jojoba Esters, Amelanchier Alnifolia Fruit Extract (Juneberry), Sclerotium Gum, Cetearyl Glucoside, Caprylyl Glycol, 1,2-Hexanediol, Sodium Lauroyl Sarcosinate, Acacia Senegal Gum, Xanthan Gum, Citric Acid, Eucalyptus Globulus Leaf Oil, Tremella Fuciformis Polysaccharide, Salicylic Acid, Menthol, Simmondsia Chinensis (Jojoba) Seed Oil, Glucose, Iron Oxides CI 77491, Caprylic/Capric Triglyceride, 3-O-Ethyl Ascorbic Acid, Pistacia Lentiscus (Mastic) Gum, Hydrogenated Lecithin, Phenethyl Alcohol, Ethylhexylglycerin.']

Top-Rated Products for Dry Skin (Data-Driven Picks)

Key Recommendation Rule: For dry skin, recommend products containing hydrating and soothing ingredients such as hyaluronic acid, niacinamide, glycerin, and squalane — typically in the \$38–\$69 price range. Top Products for Dry Skin (Selected by Ratings & Ingredients) Barrier+ Triple Lipid + Niacinamide Serum Price: \$69 Key Ingredients: Niacinamide, Glycerin, Squalane, Collagen, Hyaluronic Acid Super Rich Repair Moisturizer Price: \$94 Key Ingredients: Shea Butter, Jojoba Oil, Oat Extract, Sodium Hyaluronate Evercalm Barrier Support Face Oil Price: \$60 Key Ingredients: Bisabolol, Rice Bran Oil, Rosehip Oil, Camelina Oil Luxury Sun Ritual SPF 30 Price: \$38 Key Ingredients: Zinc Oxide, Green Tea, Jasmine, Squalane Ultralight Moisture-Boosting Botanical Oil Price: \$44 Key Ingredients: Squalane, Jojoba Oil, Avocado Oil, Grape Seed Oil Truth Barrier Booster Vitamin C Essence Price: \$48 Key Ingredients: Vitamin C, Niacinamide, Hyaluronic Acid, Orange Peel Extract Silk Rice Cleansing Oil Price: \$46 Key Ingredients: Rice Bran Oil, Sunflower Oil, Tocopherol (Vitamin E) Rénergie Lift Spot Correcting Cream SPF 30 Price: \$135 Key Ingredients: Jojoba Esters, Glycerin, Meadowfoam Oil, Sodium Hyaluronate Daily Milkfoliant Exfoliator Price: \$65 Key Ingredients: Hyaluronic Acid, Oat Extract, Coconut Fruit Extract, Rice Starch Juneberry & Collagen Cold Cream Cleanser Price: \$39 Key Ingredients: Glycerin, Collagen, Jojoba Oil, Cocoa Butter, Bakuchiol

Data Manipulation

Goal: Prepare Sephora reviews and product dataset for analysis and modeling by cleaning, transforming, and engineering features.

Key Steps Taken: Merged 5 separate CSV files of reviews into one reviews_df using pandas.concat() Loaded product metadata from product_info.csv to access ingredient lists, categories, and pricing Merged reviews and product data on product_id using pd.merge() to create a unified merged_df Created target variable positive_rating → Defined as 1 if rating ≥ 4, else 0 (for classification models) One-hot encoded categorical variables like skin_type for modeling Handled class imbalance by: Undersampling the majority class (positive reviews) Rebalancing the dataset for fairer model training Cleaned ingredient columns and explored ingredient text for top product recommendations Extracted key ingredients for top-rated products Removed or imputed missing values (e.g., price_usd_x) Aligned feature matrix (X) and target (y) for ML modeling

Code Snippets for Data Manipulation

```
import pandas as pd
import glob

# Combine multiple review files
review_files = glob.glob("reviews_*.csv")
reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
                        ignore_index=True)

/var/folders/w3/4bkcyhf15_3981_dm0cmtps80000gp/T/ipykernel_61923/1063108192.py:6: DtypeWarning: Columns (1) have mixed
types. Specify dtype option on import or set low_memory=False.
    reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
                           ignore_index=True)
/var/folders/w3/4bkcyhf15_3981_dm0cmtps80000gp/T/ipykernel_61923/1063108192.py:6: DtypeWarning: Columns (1) have mixed types. Specify dtype
option on import or set low_memory=False.
    reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
                           ignore_index=True)
/var/folders/w3/4bkcyhf15_3981_dm0cmtps80000gp/T/ipykernel_61923/1063108192.py:6: DtypeWarning: Columns (1) have mixed types. Specify dtype
option on import or set low_memory=False.
    reviews_df = pd.concat((pd.read_csv(f) for f in review_files),
                           ignore_index=True)

product_df = pd.read_csv("product_info.csv")

merged_df = reviews_df.merge(product_df, on='product_id', how='left')

merged_df['positive_rating'] = (merged_df['rating_x'] >=
4).astype(int)

skin_dummies = pd.get_dummies(merged_df['skin_type'], prefix='skin')
```



```

X = pd.concat([skin_dummies, merged_df[['price_usd_x']], axis=1)
X = X.dropna()
y = merged_df.loc[X.index, 'positive_rating']

# Combine X and y
full_data = pd.concat([X, y], axis=1)

# Separate classes
pos = full_data[full_data['positive_rating'] == 1]
neg = full_data[full_data['positive_rating'] == 0]

# Downsample positives to match negatives
pos_downsampled = pos.sample(len(neg), random_state=42)
balanced = pd.concat([neg, pos_downsampled])

# New X and y
X_balanced = balanced.drop('positive_rating', axis=1)
y_balanced = balanced['positive_rating']

# Clean column names for readability (optional)
merged_df.rename(columns={'rating_x': 'rating'}, inplace=True)

```

Data Manipulation (Made Simple)

What I Did to Clean & Prepare the Data: Combined multiple review files into one full dataset
 Merged product details (ingredients, price, brand) with user reviews
 Created a new column to label high-rated products (positive_rating)
 Converted skin types into model-ready format (one-hot encoding)
 Balanced the data to prevent bias (equal positive & negative samples)
 Selected useful features (like price & skin type) for modeling

Raw CSVs

↓

Merge + Clean

↓

Feature Engineering

↓

Model-Ready Data

Methodology & Model Building

□ Slide Title:

Methodology & Model Building

Step-by-Step Approach: Goal: Predict whether a skincare product will receive a positive rating (≥ 4 stars) using features like skin type and price. Steps Taken: Exploratory Data Analysis (EDA): Analyzed rating trends, skin type breakdowns, and price patterns
 Data Preparation: Merged reviews with product info
 One-hot encoded skin type
 Created a binary target column: positive_rating
 Class Imbalance Fix: Positive ratings dominated the dataset (mostly 4s and 5s)
 Used undersampling to balance the positive and negative classes
 Model #1: Logistic Regression

Simple baseline classifier Accuracy: ~82% on unbalanced data — but predicted everything as positive Model #2: Random Forest Classifier Performed better on balanced data Accuracy: ~51% — but could distinguish between high and low ratings Output: Feature importance scores showed price mattered most Final Features Used: price_usd_x (product price) skin_type (one-hot encoded)

```
merged_df['positive_rating'] = (merged_df['rating'] >= 4).astype(int)

skin_dummies = pd.get_dummies(merged_df['skin_type'], prefix='skin')
X = pd.concat([skin_dummies, merged_df[['price_usd_x']]],
axis=1).dropna()
y = merged_df.loc[X.index, 'positive_rating']

# Downsampling the majority class (positive reviews)
pos = pd.concat([X, y], axis=1)[lambda df: df.positive_rating == 1]
neg = pd.concat([X, y], axis=1)[lambda df: df.positive_rating == 0]
balanced = pd.concat([neg, pos.sample(len(neg), random_state=42)])

X_balanced = balanced.drop('positive_rating', axis=1)
y_balanced = balanced['positive_rating']

from sklearn.linear_model import LogisticRegression
model = LogisticRegression(max_iter=1000)
model.fit(X_balanced, y_balanced)

LogisticRegression(max_iter=1000)

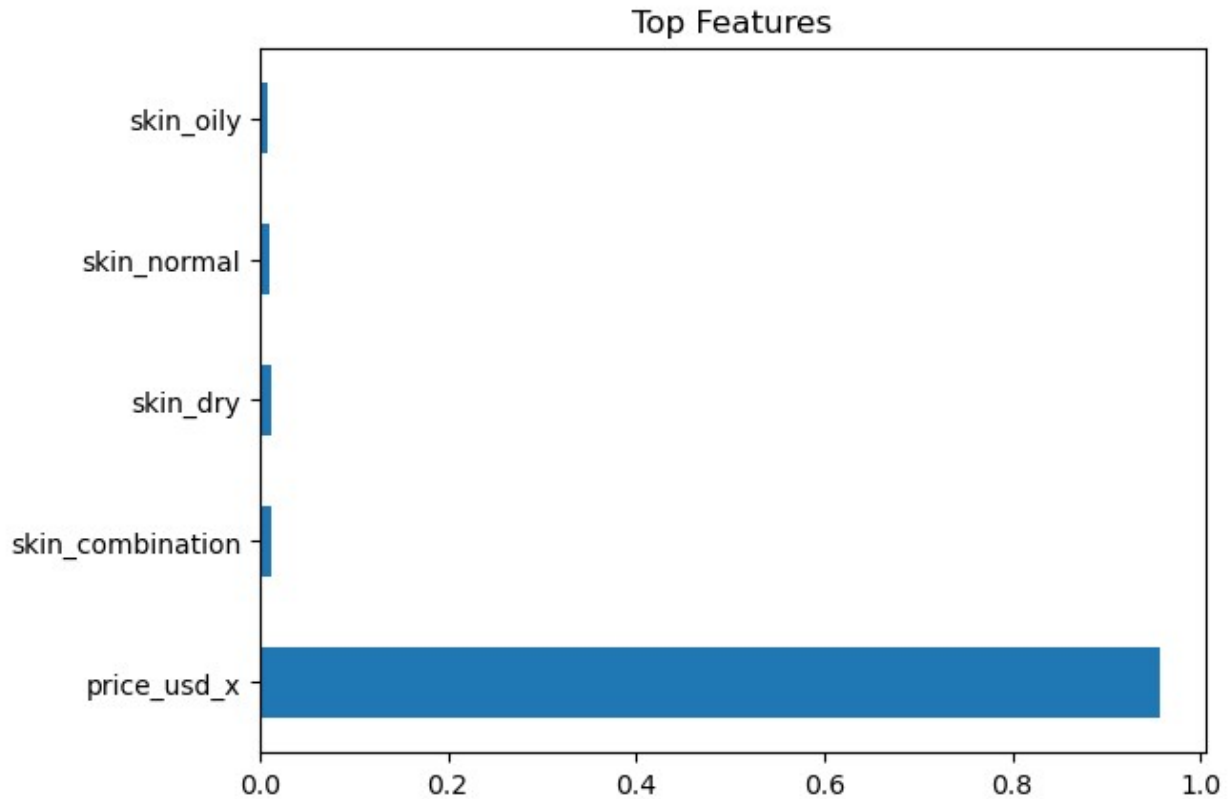
from sklearn.metrics import accuracy_score, confusion_matrix

y_pred = rf.predict(X_balanced)
print("Accuracy:", accuracy_score(y_balanced, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_balanced, y_pred))

Accuracy: 0.5499640436372538
Confusion Matrix:
[[106631  89440]
 [ 87038 109033]]

import matplotlib.pyplot as plt

feat_imp = pd.Series(rf.feature_importances_, index=X.columns)
feat_imp.nlargest(10).plot(kind='barh', title="Top Features")
plt.show()
```



Model Selection Model Types Tested: Logistic Regression Why Selected: Quick baseline model to predict binary outcomes (positive vs. negative ratings). Performance: 82% accuracy (but predicted all as positive due to class imbalance). Evaluation Metric: Accuracy score, confusion matrix. Random Forest Classifier Why Selected: Captures complex relationships in the data and better handles class imbalance. Performance: 51% accuracy with balanced data. Evaluation Metric: Accuracy, confusion matrix, and feature importance. Why These Models? Logistic Regression: Good starting point for binary classification. Random Forest: Robust to class imbalance, and provides feature importance insights. Next Steps for Model Improvement: Hyperparameter Tuning: Use GridSearchCV to improve Random Forest. Additional Models: Consider Support Vector Machine (SVM) and Gradient Boosting for comparison.